

# Chapter 4

# Network

# Layer:

# Data Plane

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2020  
J.F Kurose and K.W. Ross, All Rights Reserved



## Computer Networking: A Top-Down Approach

8<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson, 2020

# Network layer: our goals

- understand principles behind network layer services, focusing on **data plane**:
  - network layer service models
  - forwarding versus routing
  - how a router works
  - addressing
  - generalized forwarding
- IP protocol
- Network address translation (NAT)

# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- forwarding plane (data plane)
- routing (control plane)

## 4.2 What's inside a router

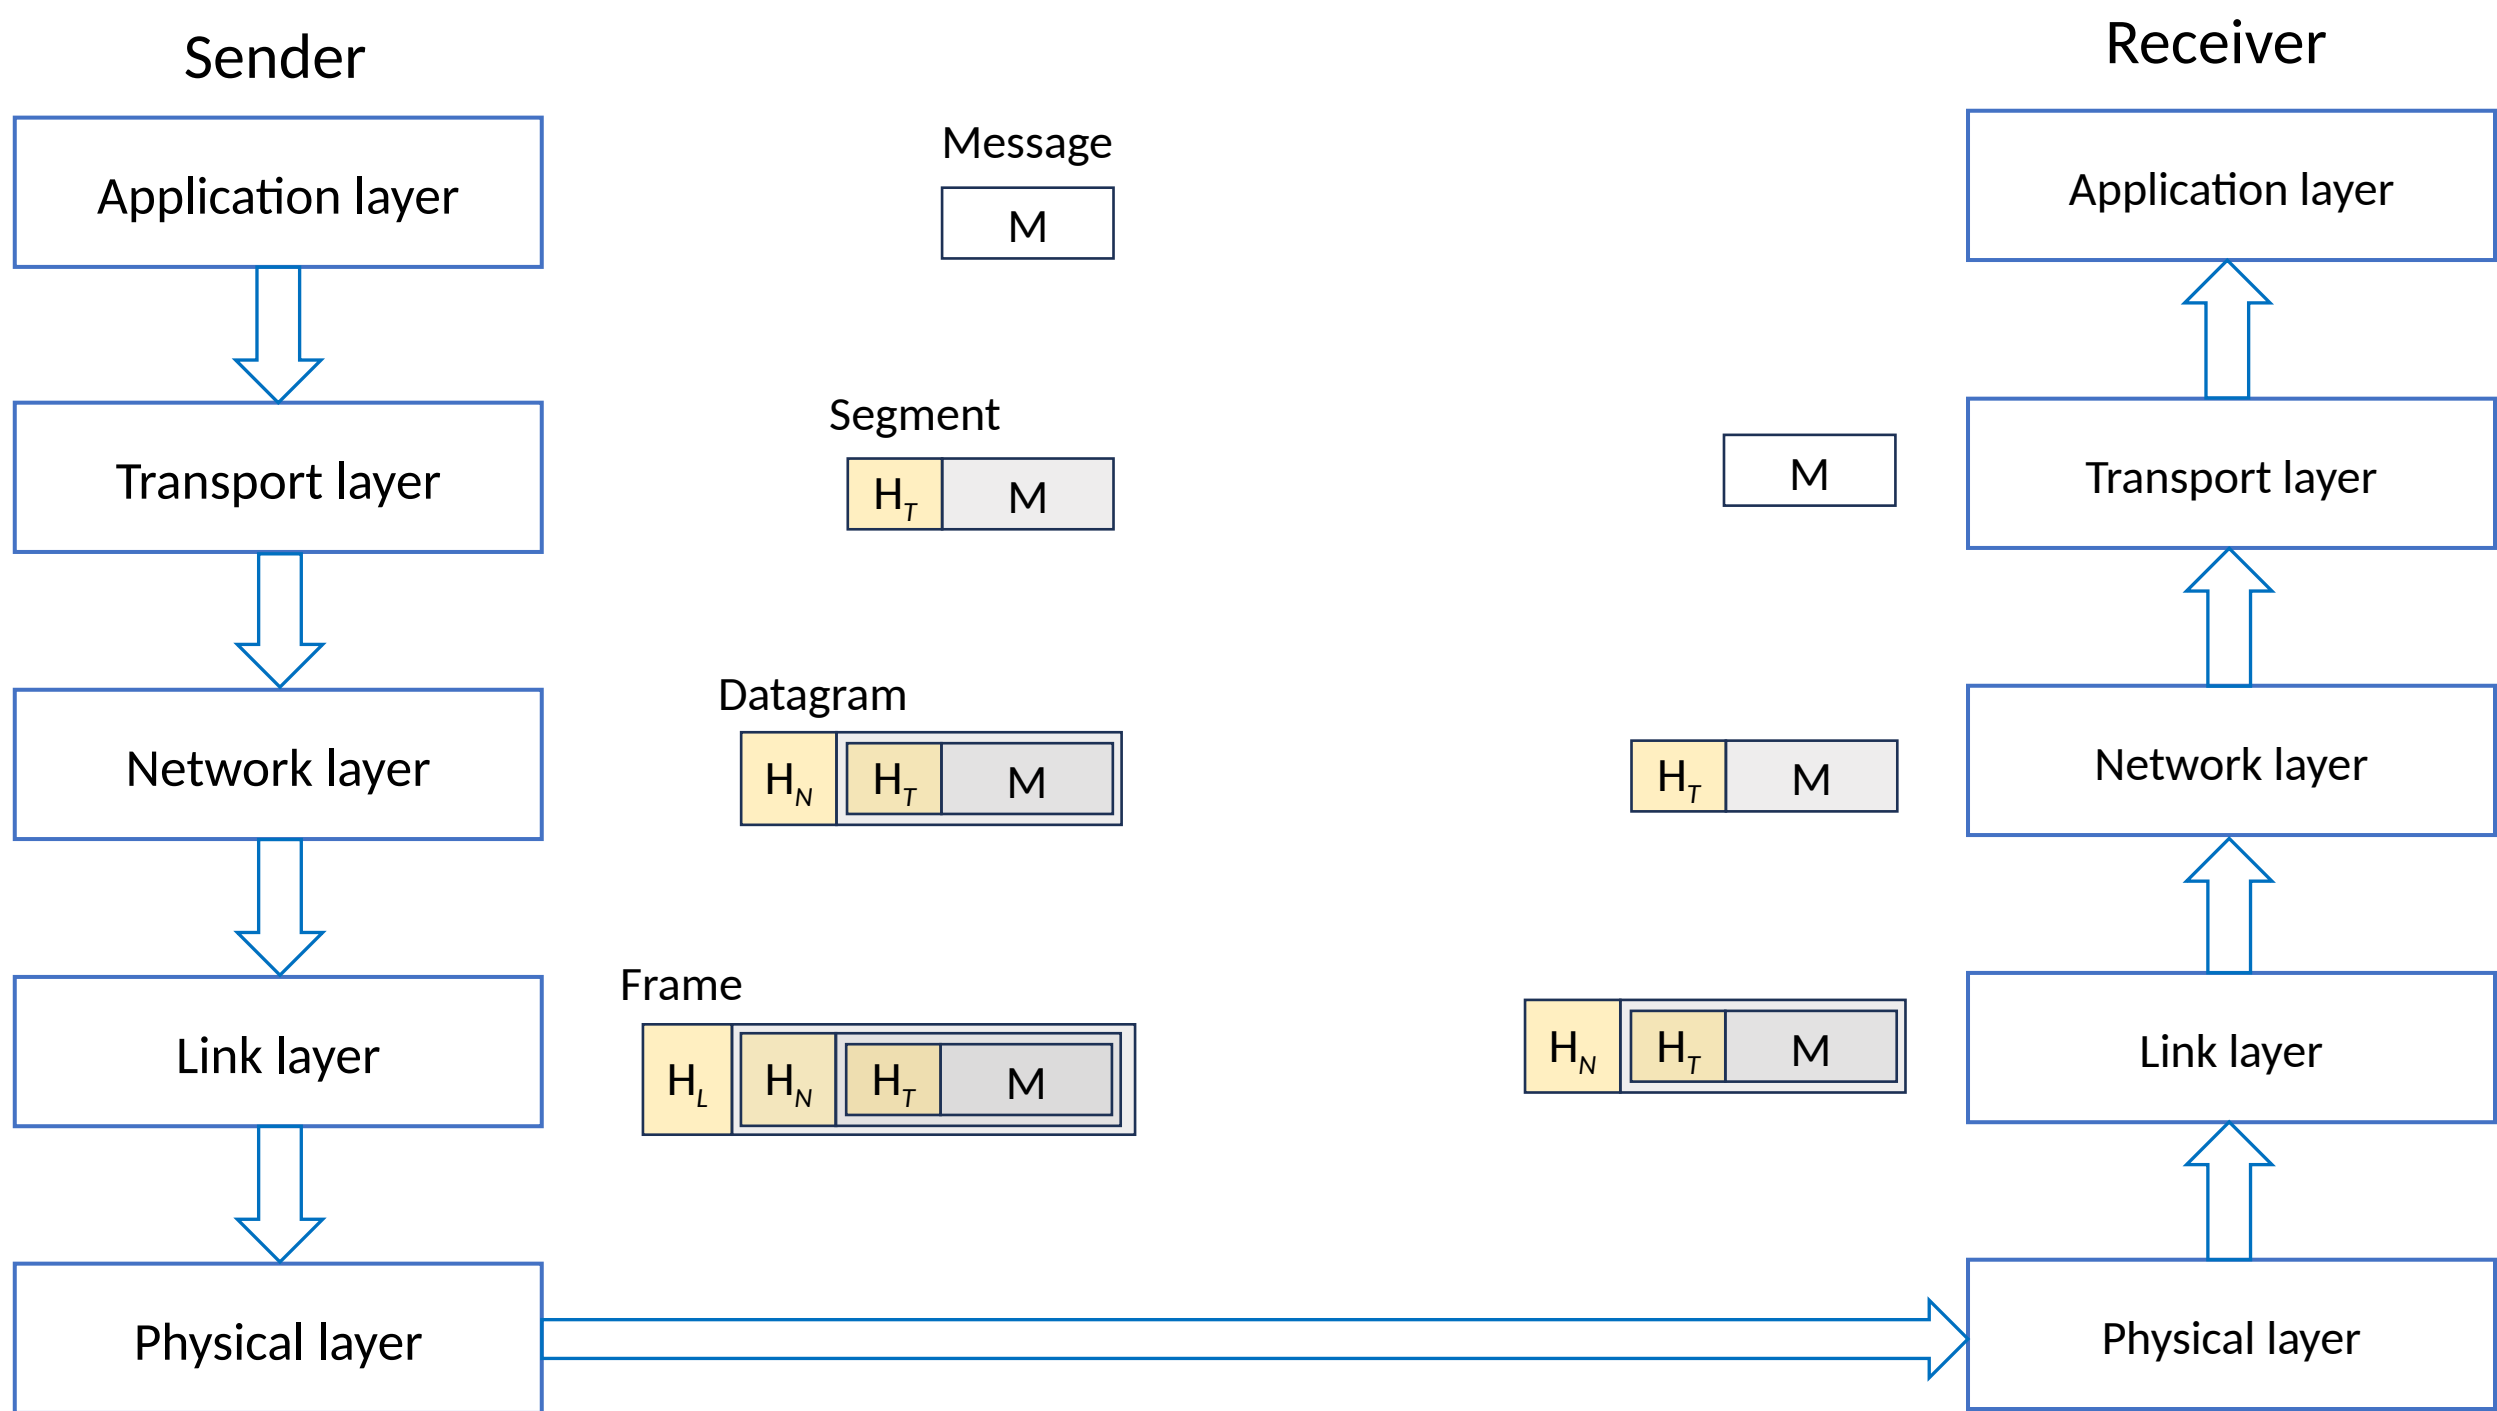
- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6

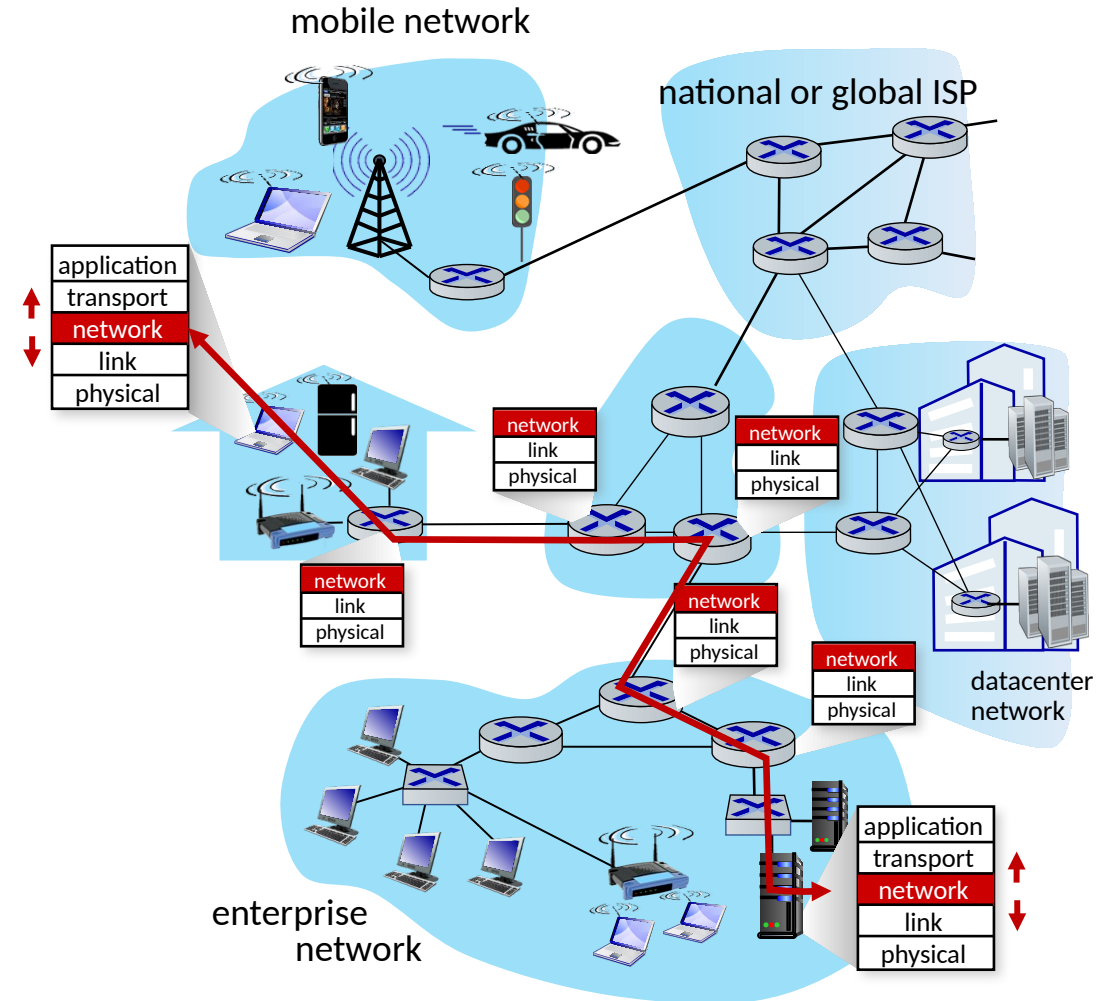
## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action



# Network-layer services and protocols

- transport segment from sending to receiving host
  - **sender:** encapsulates segments into datagrams, passes to link layer
  - **receiver:** delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
  - examines header fields in all IP datagrams passing through it
  - moves datagrams from input ports to output ports to transfer datagrams along end-end path



# Two key network-layer functions

## network-layer functions:

- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*
- *forwarding*: move packets from a router's input link to appropriate router output link

## analogy: taking a trip

- *routing*: process of planning trip from source to destination



routing


- *forwarding*: process of getting through single interchange

# Network layer: forwarding plane, control plane

## Forwarding (“data plane”):

- *local*, per-router function
- determines how datagram arriving on router input port is **forwarded** to router output port
- forwarding table

values in arriving  
packet header

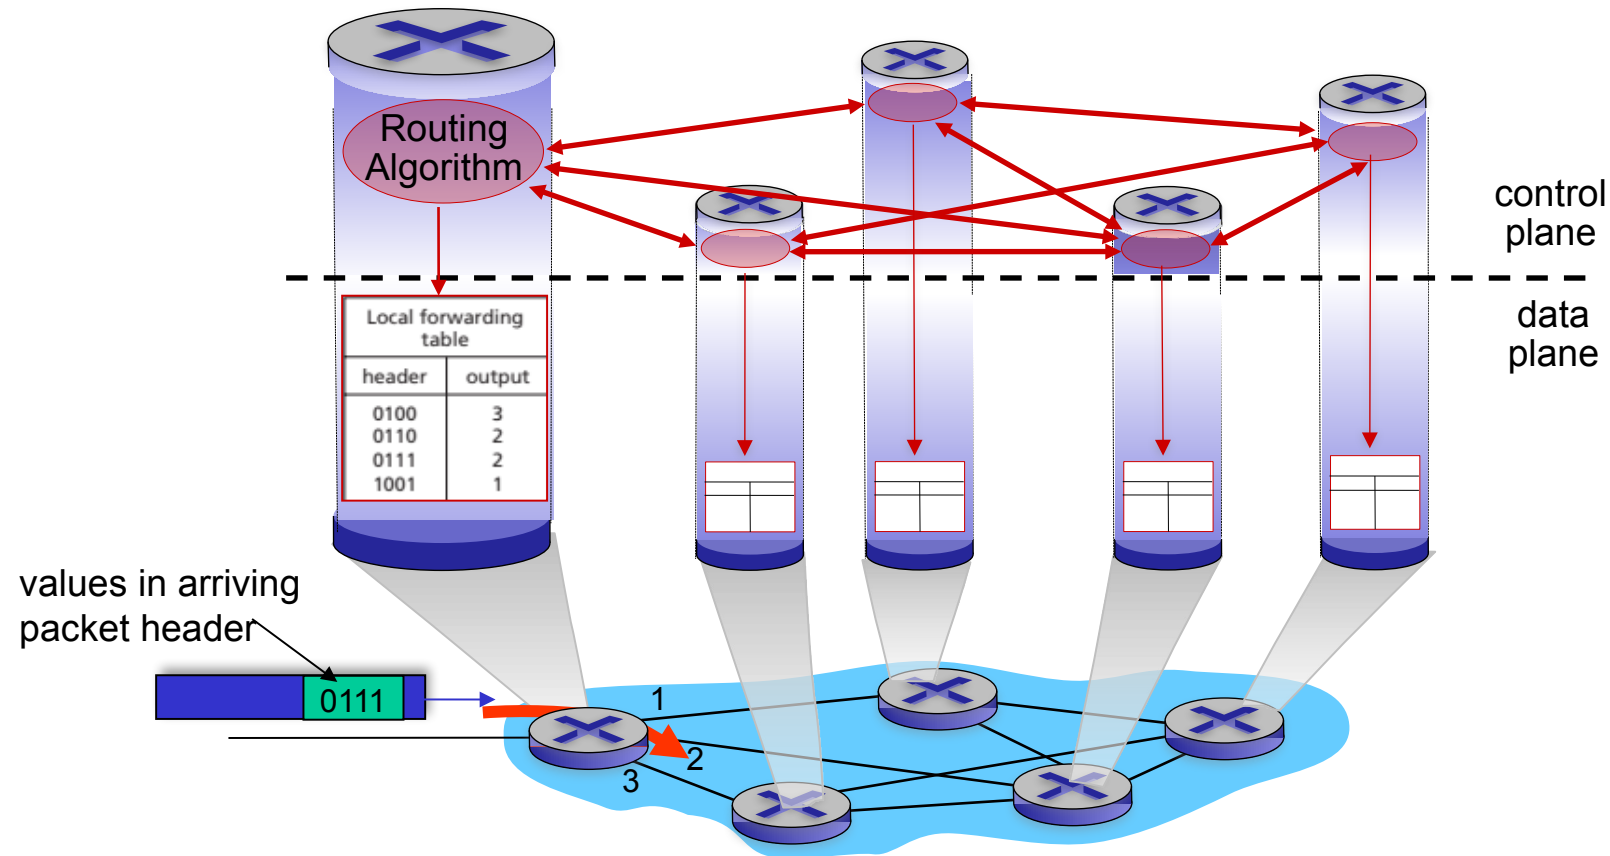


## Routing (“control plane”):

- *network-wide* logic
- determines how datagram is **routed** among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *Distributed routing*: implemented in routers
  - *Centralized routing (aka. software-defined networking)*: implemented in remote servers

# Distributed routing

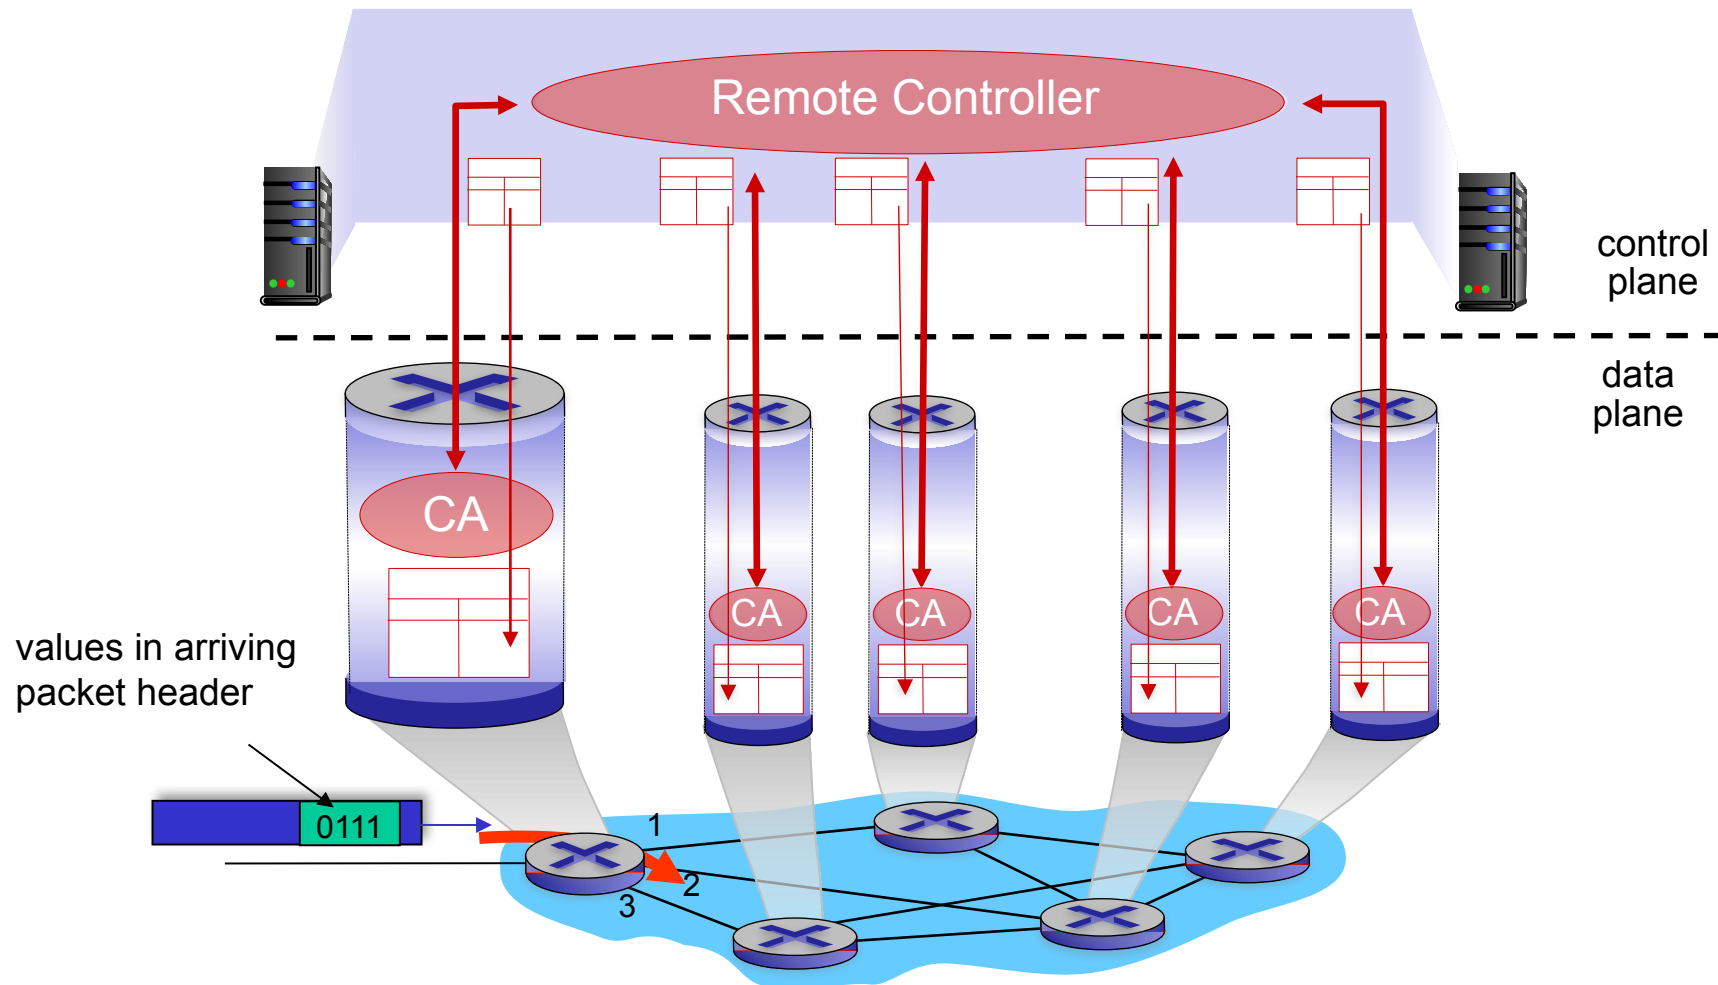
Individual routing algorithms *in each and every router*





# Centralized routing: Software-Defined Networking (SDN)

Remote controller computes forwarding tables in routers



# IP network-layer service model

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no

Internet “best effort” service model

**No** guarantees on:

- i. successful datagram delivery (r.d.t.)
- ii. order of delivery (r.d.t.)
- iii. bandwidth available to end-end flow

# Network layer: “forwarding plane” roadmap

## 4.1 Network layer: overview

- forwarding plane
- control plane

## 4.2 What’s inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 IP: the Internet Protocol

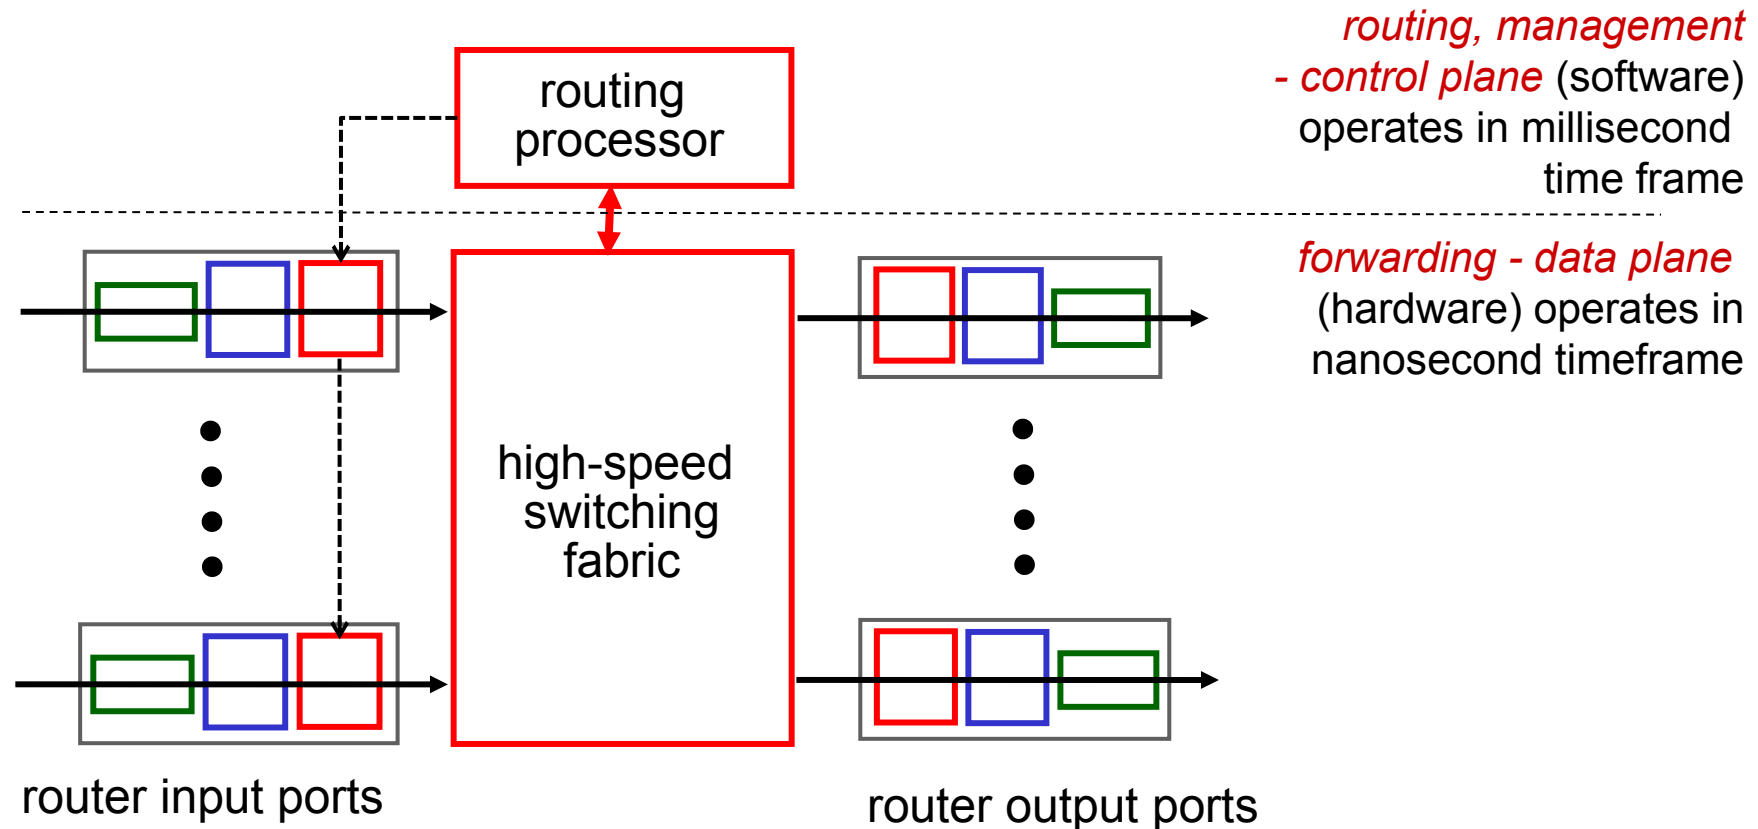
- datagram format
- addressing
- network address translation
- IPv6

## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action

# Router architecture overview

high-level view of generic router architecture:



# Destination-based forwarding

<i>forwarding table</i>	
Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010000 00000100	n
11001000 00010111 00010000 00000111	
11001000 00010111 00011000 11111111	
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	3
otherwise	2

**Q:** but what happens if ranges don't divide up so nicely?

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000    00010111    00010***    *****	0
11001000    00010111    00011000    *****	1
11001000    00010111    00011***    *****	2
otherwise	3

examples:    11001000    00010111    00010110    10100001    which interface?

11001000    00010111    00011000    10101010    which interface?

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 match! 1 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?  
11001000 00010111 00011000 10101010 which interface?

# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

match!

examples:

11001000 00010111 00010110 10100001	which interface?
11001000 00010111 00011000 10101010	which interface?



# Longest prefix matching

## longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

match!

examples:

11001000 00010111 00010110 10100001	which interface?
11001000 00010111 00011000 10101010	which interface?

# Longest prefix matching

## longest prefix match

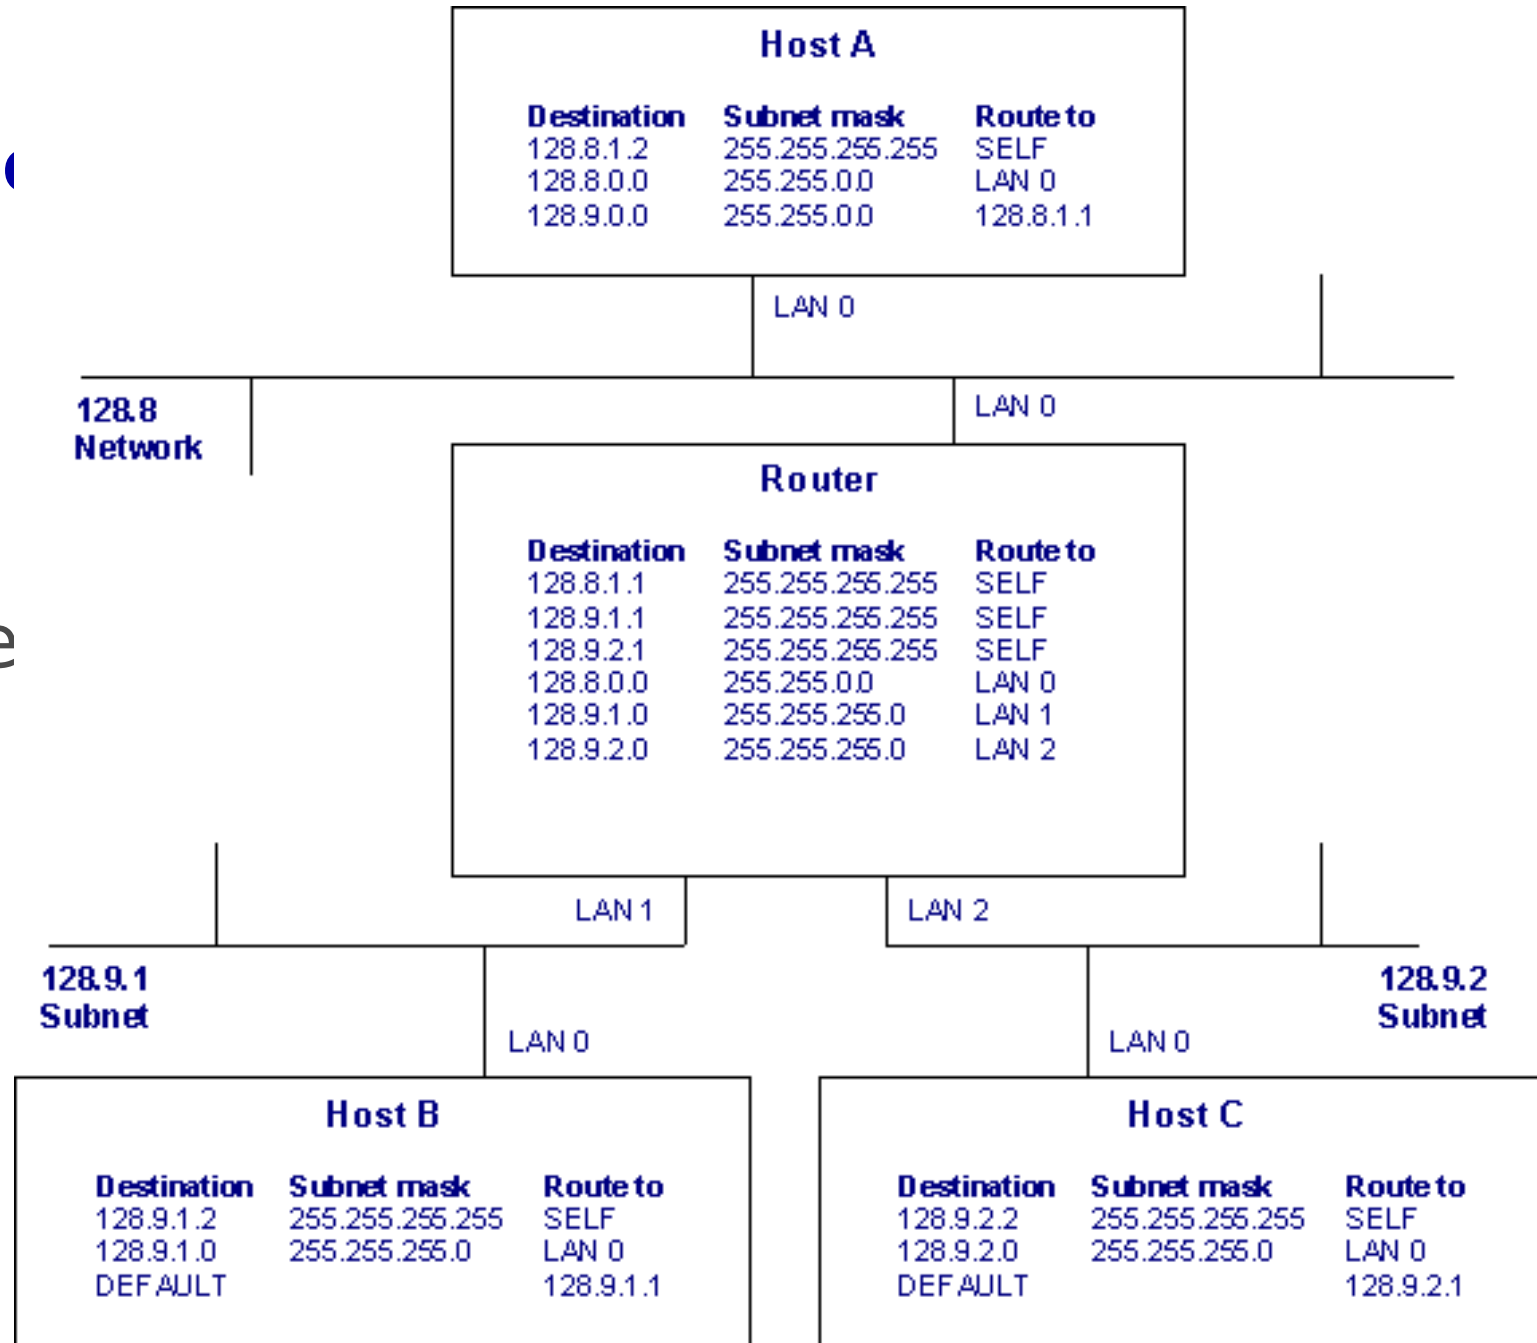
when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

new entry

Destination Address Range				Link interface
11001000	00010111	00010***	*****	0
11001000	00010111	00011000	*****	1
11001000	00010111	00011***	*****	2
<b>11001000</b>	<b>00010111</b>	<b>00010000</b>	<b>000001**</b>	<b>3</b>
otherwise				3




# Destination-based forwarding example

- Host A sends a packet to **128.9.2.2** (Host C)
- The router has three addresses
  - one IP address per interface



# Switching fabrics

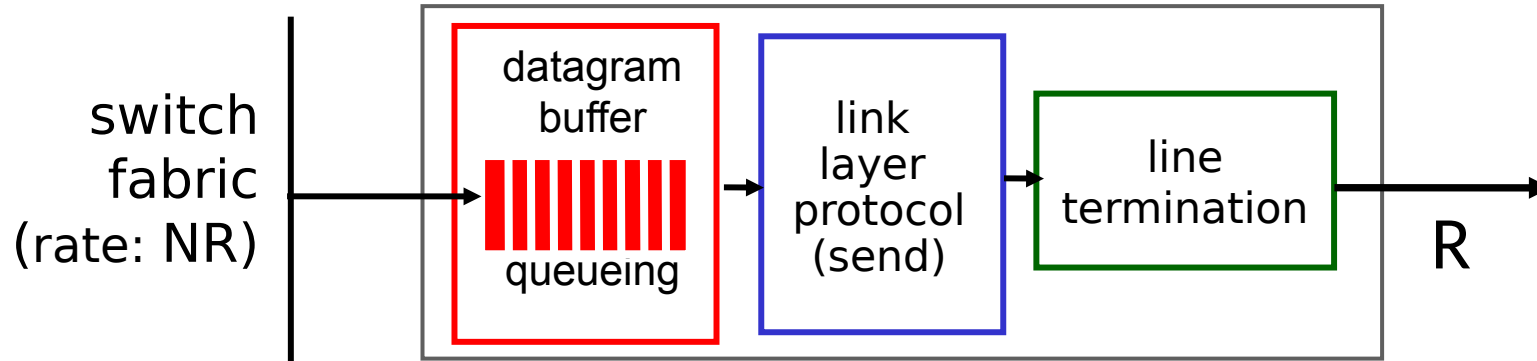
- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transfer from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable
- three major types of switching fabrics:



# Output port queuing



This is a really important slide



- **Buffering** required when datagrams arrive from fabric faster than link transmission rate. **Drop policy:** which datagrams to drop if no free buffers?



Datagrams can be lost due to congestion, lack of buffers

- **Scheduling discipline** chooses among queued datagrams for transmission




Priority scheduling – who gets best performance, network neutrality

# Packet scheduling

**packet scheduling:** deciding which packet to send next on link:

1. First come, first served (FCFS)
2. Priority
3. Round robin
4. Weighted fair queueing

Abstraction: queue




# Packet Scheduling: First-come-first-served (FCFS)

## FCFS:

- packets transmitted in the same order they arrive the queue
- also known as: First-in-first-out (FIFO)
- real world examples?



Abstraction: queue



# Scheduling policies: Priority scheduling

## *Priority scheduling:*

- arriving traffic classified, queued by class
  - any header fields can be used for classification
- send packet from highest priority queue that has buffered packets
  - FCFS within priority class






# Scheduling policies: Round-robin scheduling

## *Round Robin (RR) scheduling:*

- arriving traffic classified, queued by class
  - any header fields can be used for classification
- server cyclically, repeatedly scans class queues, sending one complete packet from each class (if available) in turn




# Scheduling policies: weighted fair queueing

## *Weighted Fair Queuing (WFQ):*

- generalized Round Robin
- each class,  $i$ , has weight,  $w_i$ , and gets weighted amount of service in each cycle:

$$\frac{w_i}{\sum_j w_j}$$

- minimum bandwidth guarantee (per-traffic-class)



# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- data plane
- control plane

## 4.2 What's inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 The Internet Protocol (IP)


- datagram format
- addressing
- network address translation
- IPv6

## 4.4 Generalized Forwarding, SDN

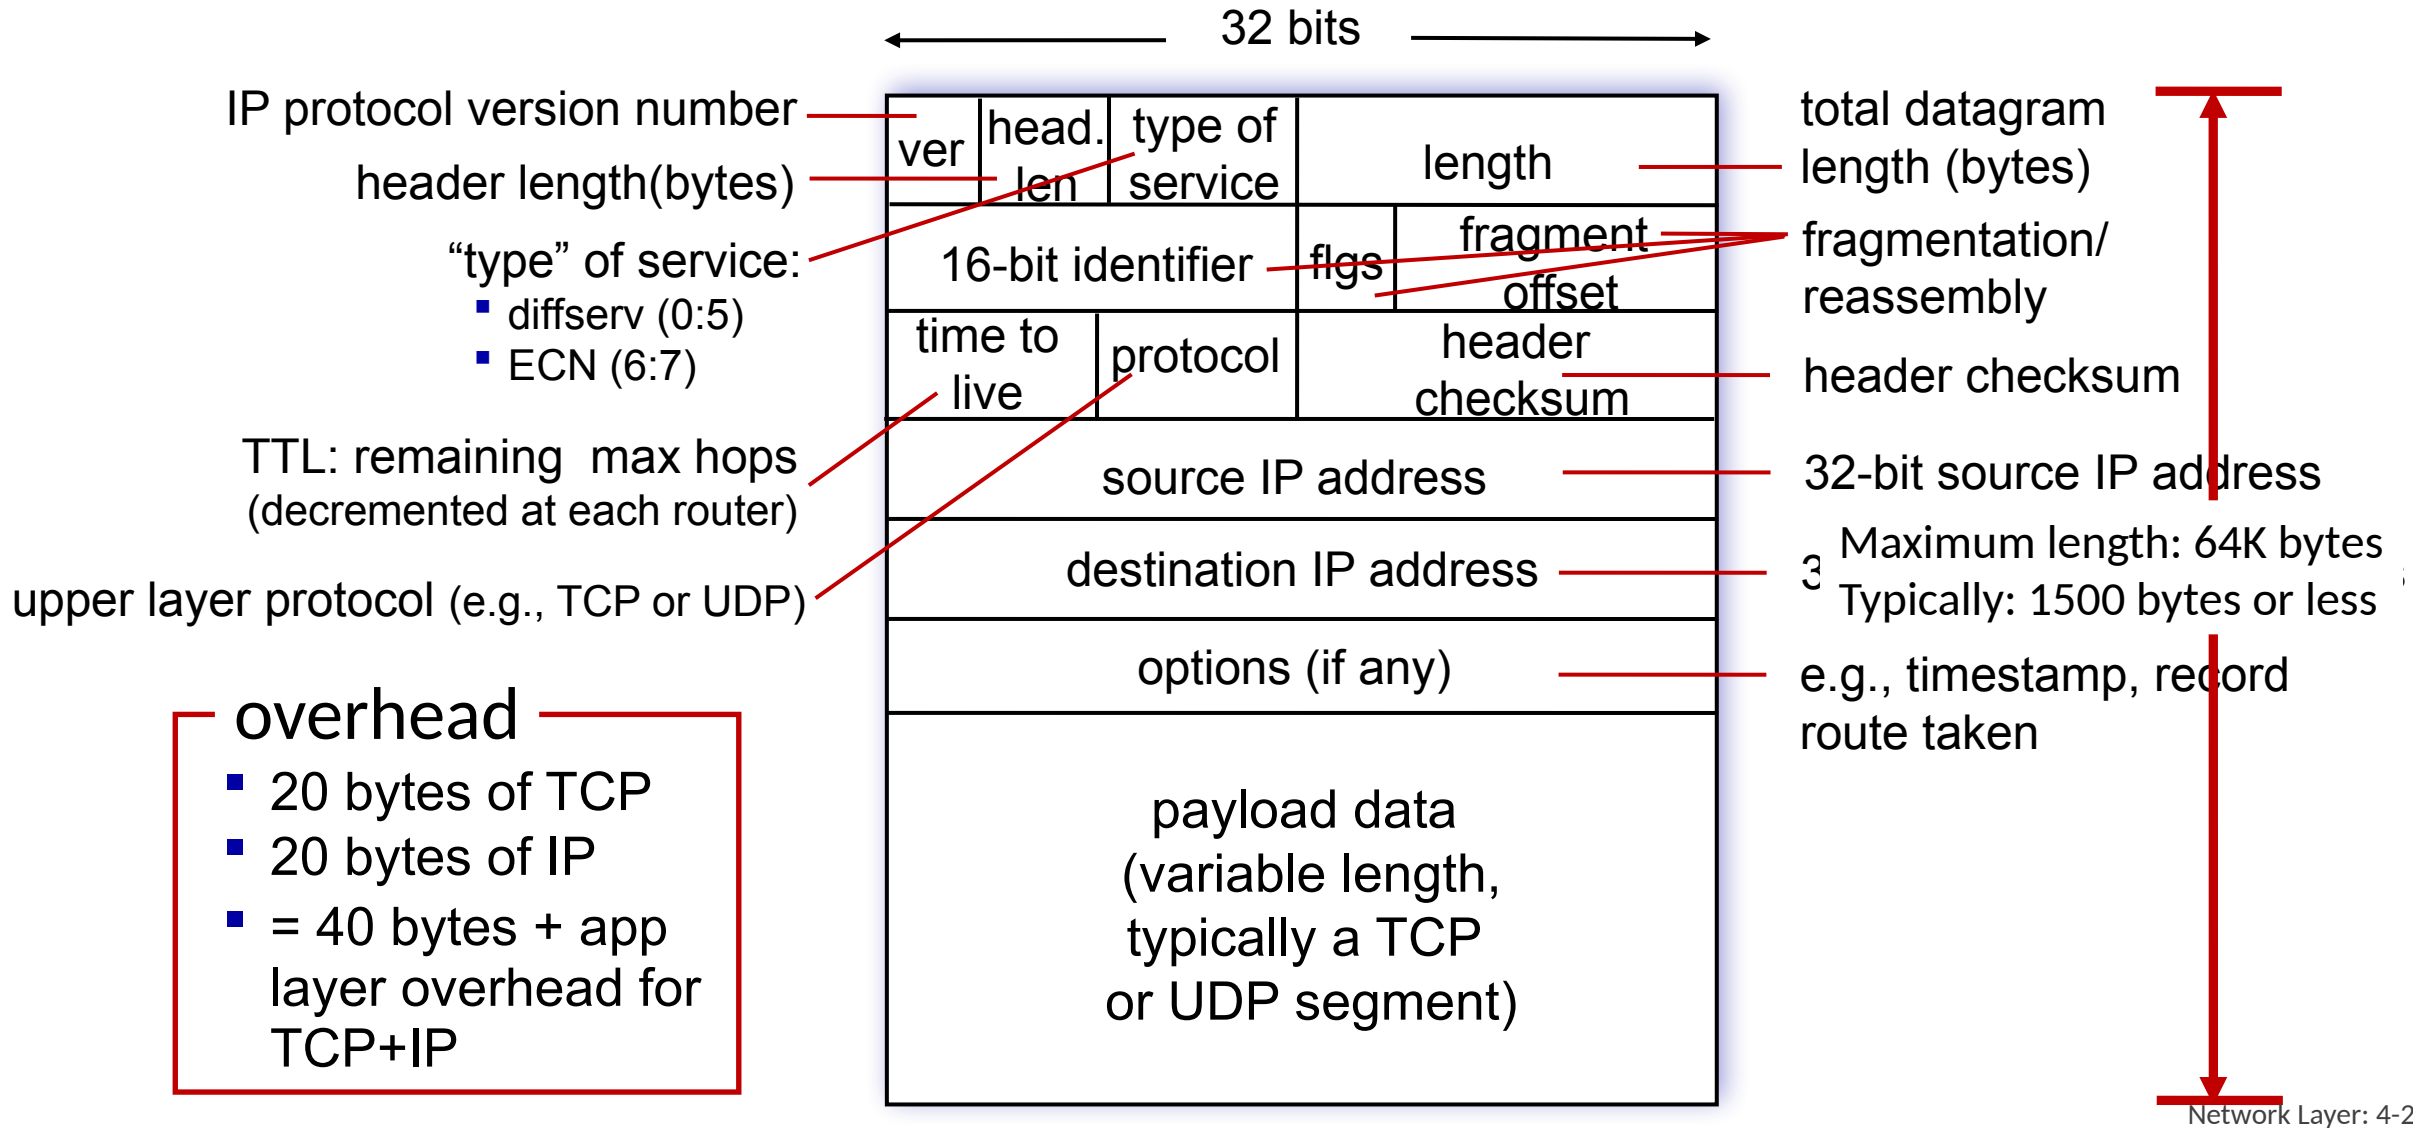
- match + action
- OpenFlow: match + action in action

# Network Layer: Internet

host, router network layer functions:



# IPv4 Datagram format



# Maximum payload size: MTU vs. MSS

Maximum payload size:

- **Frames:**

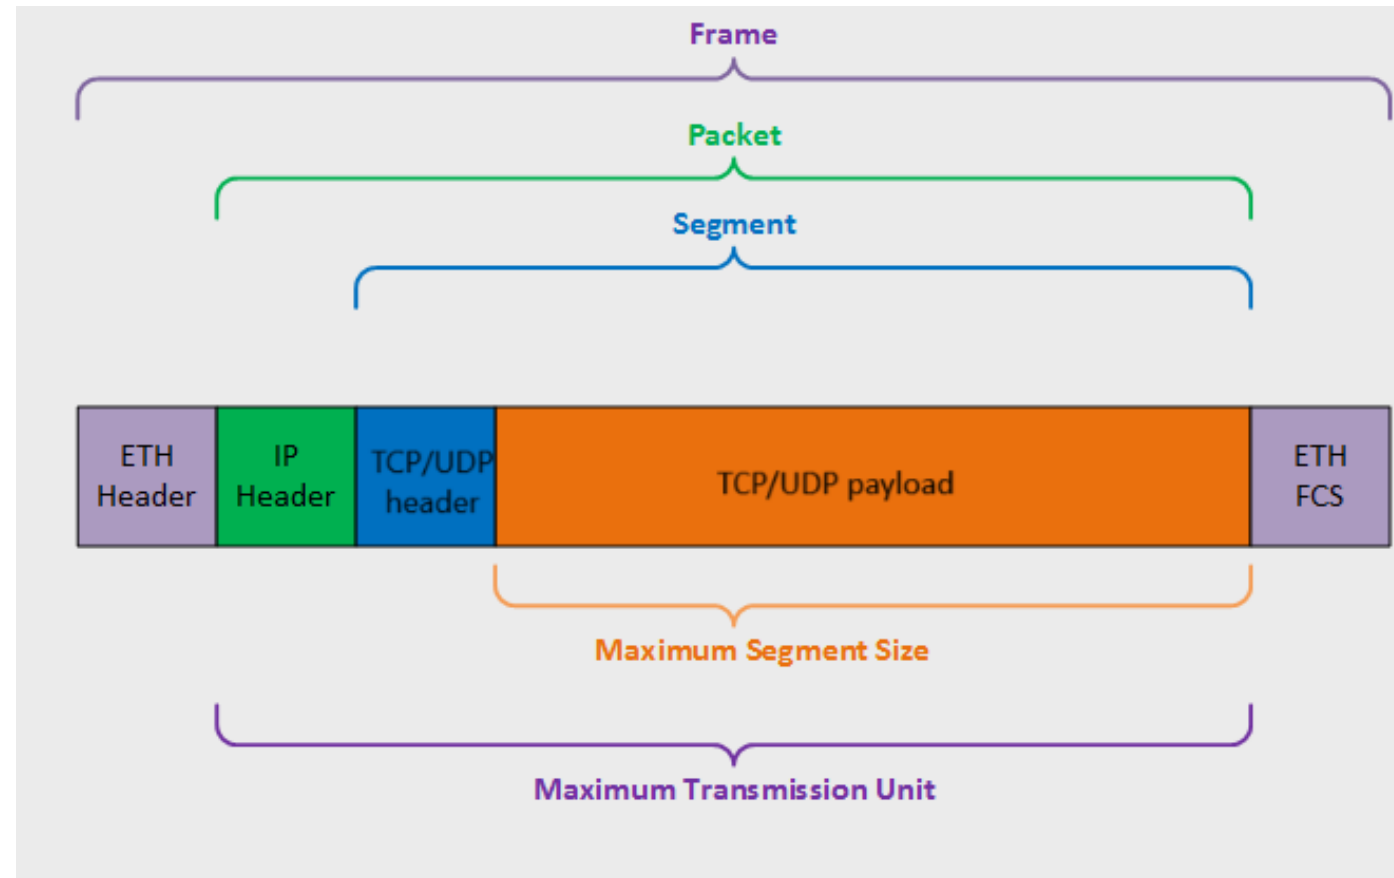
Maximum transfer unit (MTU)

- e.g. Ethernet: 1500 bytes

- **Segments:**


Maximum segment size (MSS)

- $MSS = MTU - 40$   
=  $1500 - 40 = 1460$
- 40 = IP header size + TCP header size

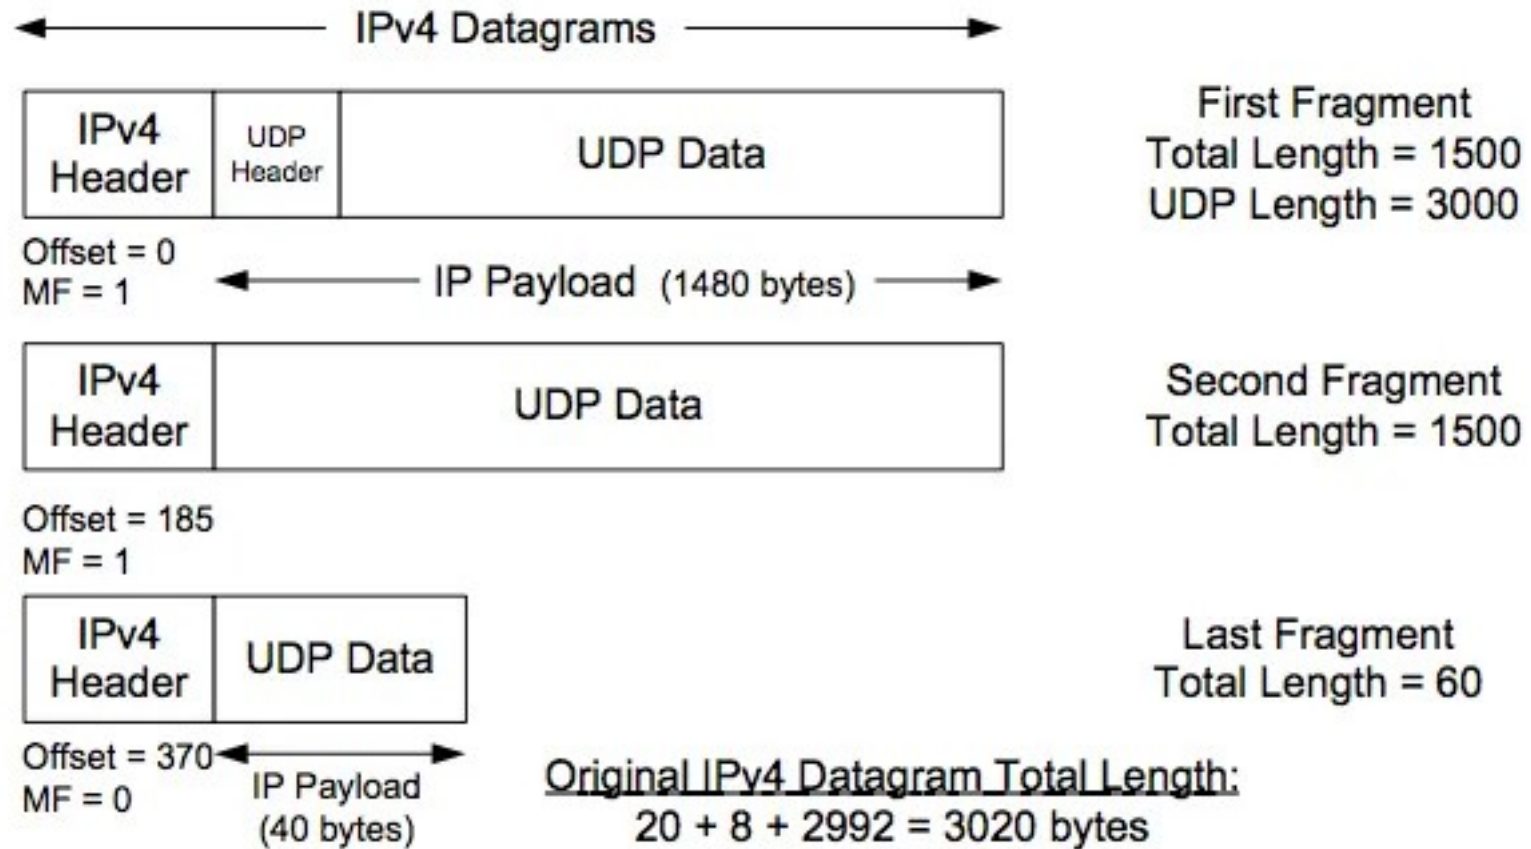


# IP fragmentation

- network links have MTU (maximum transfer unit) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at *destination*
  - IP header bits used to identify, order related fragments



# IP fragmentation





# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- data plane
- control plane

## 4.2 What's inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 The Internet Protocol (IP)


- datagram format
- addressing and subnets
- network address translation
- IPv6

## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action

# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - routers have **multiple interfaces**
  - hosts typically have one or two interfaces (e.g., wired Ethernet, wireless 802.11)
  - Each interface has **its own IP-address**



dotted-decimal IP address notation:

223.1.1.1 =  $\underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$

Network Layer: 4-35


# IP addressing: introduction

**Q:** how are interfaces actually connected?

**A:** we'll learn about that in chapters 6, 7

*For now:* don't need to worry about how one interface is connected to another (with no intervening router)


**A:** wired Ethernet interfaces connected by Ethernet switches



**A:** wireless WiFi interfaces connected by WiFi base station

# Subnets


- *What's a subnet ?*
  - device interfaces that can physically reach each other **without passing through an intervening router**




network consisting of 3 subnets

# Subnets

- **CIDR: Classless InterDomain Routing**
- IP addresses have structure: **a.b.c.d/x**
  - **prefix: subnet part/network part:** devices in subnet have same **x** most significant bits (msb.)
  - **host part:** remaining least significant bits



200.23.16.0/**23**



network consisting of 3 subnets

# Subnet mask vs. prefix **a.b.c.d/x**

- A subnet mask is a **bitmask** when applied by a **bitwise AND** operation to any **IP address** in the network, yields the routing **prefix** / subnet address

prefix = IP address && bitmask

- 198.51.100.20/24 has the subnet mask 255.255.255.0 (24 msb. are 1's)

**11000110 00110011 01100100** 00010100

☾ 198.51.100.20 (IP address)

**bitwise AND**

**11111111 11111111 11111111** 00000000

☾ 255.255.255.0 (subnet mask)


---

**11000110 00110011 01100100** 00000000

☾ 198.51.100.20 (subnet address)

# Subnets

- where are the subnets?
- what are the /24 subnet addresses?



# How to get IP addresses?

That's **two** questions:

1. How does an *organization* get IP addresses for itself (network part of address)
2. How does a *host* get an IP address within its network (host part of address)?



# 1. How gets an organization IP addresses?

**A:** gets allocated portion of its provider ISP's address space

ISP's block      11001000 00010111 00010000 00000000    200.23.16.0/**20**

ISP can then allocate out its address space in 8 blocks:

Organization 0	<u>11001000 00010111 0001</u> <b>000</b> 0	00000000	200.23.16.0/ <b>23</b>
Organization 1	<u>11001000 00010111 0001</u> <b>001</b> 0	00000000	200.23.18.0/ <b>23</b>
Organization 2	<u>11001000 00010111 0001</u> <b>010</b> 0	00000000	200.23.20.0/ <b>23</b>
...	.....	....	....
Organization 7	<u>11001000 00010111 0001</u> <b>111</b> 0	00000000	200.23.30.0/ <b>23</b>

## 2. How gets a host an IP address?

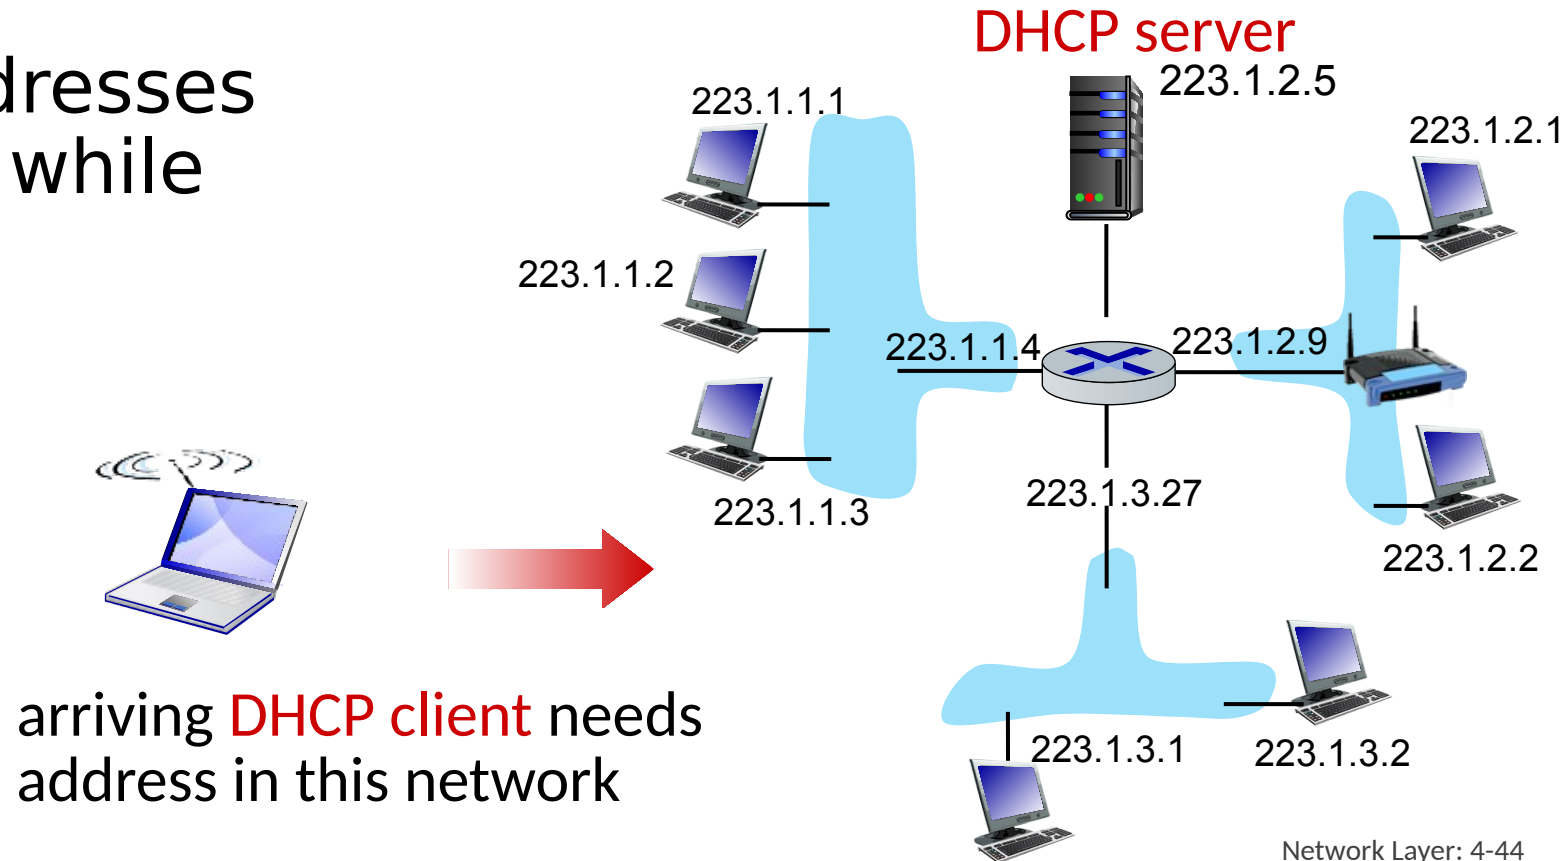
1. Manually by system admin in config file (e.g., /etc/rc.config in UNIX)
2. Dynamically from a server: **DHCP: Dynamic Host Configuration Protocol**
  - “plug-and-play”

# DHCP client-server scenario

**goal:** host *dynamically* obtains IP address from network server when it “joins” network

- allows reuse of addresses (only hold address while connected/on)

Typically, DHCP server will be co-located in router, serving all subnets to which router is attached



# DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP discover

Broadcast: is there a  
DHCP server out  
there?

Arriving client



DHCP offer

Broadcast: I'm a DHCP  
server! Here's an IP  
address you can use

DHCP request

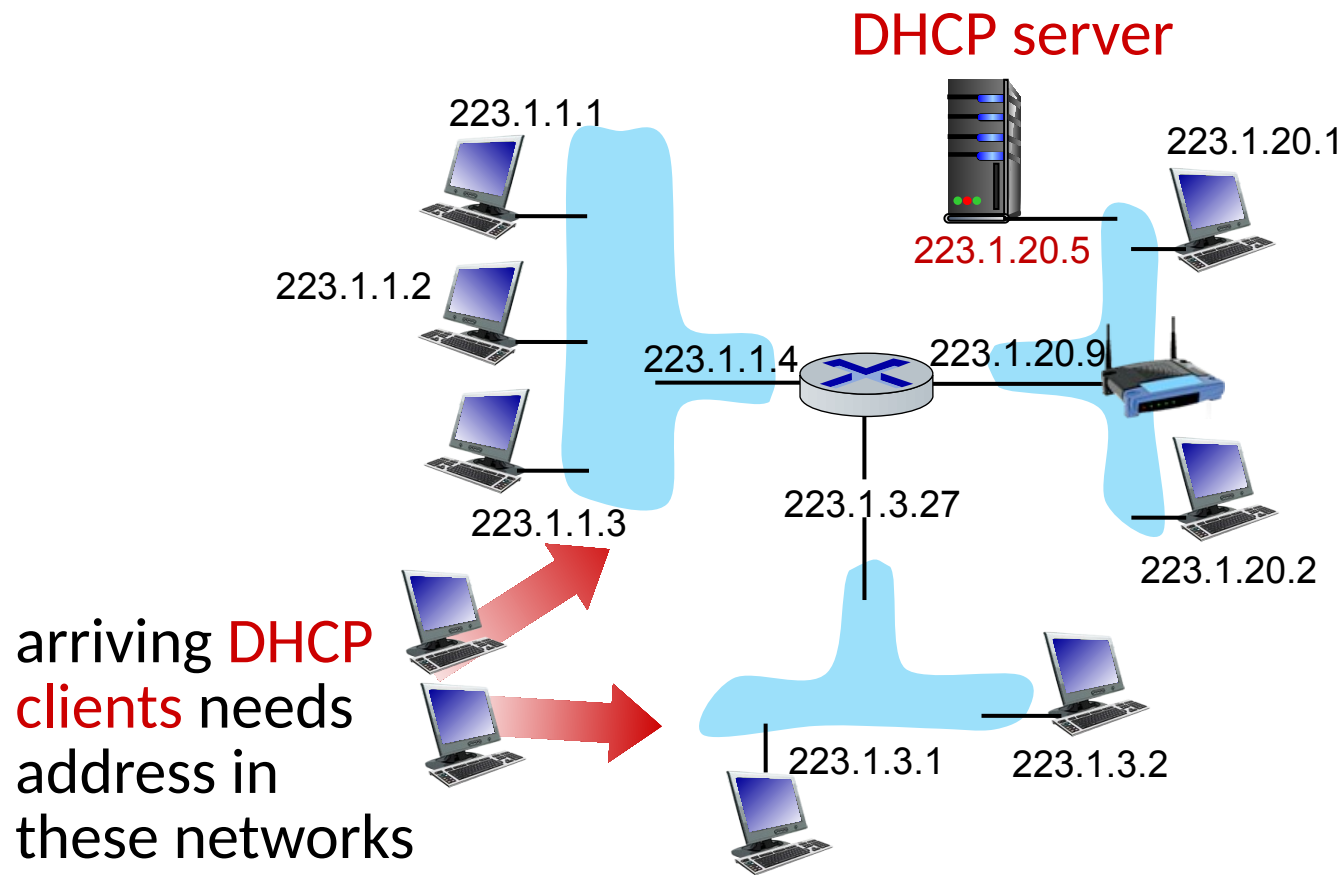
Broadcast: OK. I would  
like to use this IP  
address!

DHCP ACK

Broadcast: OK. You've  
got that IP address!

The two steps above can  
be skipped "if a client  
remembers and wishes to  
reuse a previously  
allocated network address"  
[RFC 2131]

# DHCP router relay agent



DHCP server located outside router serving all subnets to which router is attached:

- Router configured as a DHCP relay agent (ip helper-address 223.1.20.5)
- Router forward DHCP requests and replies between client and DHCP server
- Router sets the gateway IP address (223.1.1.4) in **giaddr** field of the DHCP packet
- This allows DHCP server to identify which subnet the request originated

# DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

- address of gateway router for client
- name and IP address of DNS sever
- subnet mask (indicating network versus host portion of address)
- IP address lease time

# IP addressing: last words ...

**Q:** how does an ISP get block of addresses?

**A:** ICANN: Internet Corporation for Assigned Names and Numbers  
<http://www.icann.org/>

- allocates IP addresses, through 5 regional registries (RRs) (who may then allocate to local registries)
- manages DNS root zone, including delegation of individual TLD (.com, .edu , ...) management

**Q:** are there enough 32-bit IP addresses?

- ICANN allocated last chunk of IPv4 addresses to RRs in 2011
- NAT (next) helps IPv4 address space exhaustion

"Who the hell knew how much address space we needed?" Vint Cerf (reflecting on decision to make IPv4 address 32 bits long)

# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- data plane
- control plane

## 4.2 What’s inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 The Internet Protocol (IP)

- datagram format
- addressing
- NAT: network address translation
- IPv6


## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action



# NAT: network address translation

**NAT:** All devices in local network share just **one** IPv4 address as far as outside world is concerned



*All* datagrams *leaving* local network have *same* source NAT IP address: 138.76.29.7, but *different* source port numbers

Datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

# NAT: network address translation

**implementation:** NAT router must (transparently):

- **outgoing datagrams: replace** (source IP address, port number) of every outgoing datagram to (public NAT IP address, new port number)
  - remote clients/servers will respond using (public NAT IP address, new port number) as destination address
- **remember (in NAT translation table)** every (source IP address, port number) to (public NAT IP address, new port number) translation pair
- **incoming datagrams: replace** (public NAT IP address, new port number) in destination fields of every incoming datagram with corresponding (source IP address, port number) stored in NAT table

# Clients in a local network (LAN)

## Network address translation (NAT)

Private network side		Public network side	
Source IP	Source Port	Target IP	NAT Port
10.0.0.1	837	3.3.3.3	267

## Remote server

### NAT router

1. changes packet's source address from 10.0.0.1:837 to 4.2.1.5:267
2. updates table



10.0.0.1

source IP: 10.0.0.1  
source port: 837  
target IP: 3.3.3.3  
target port: 80

req 1

source IP: 3.3.3.3  
source port: 80  
target IP: 10.0.0.1  
target port: 837

resp 1

Default gateway  
IP address:  
10.0.0.9

Public  
IP address:  
4.2.1.5


Gateway-router  
with  
NAT

source IP: 4.2.1.5  
source port: 267  
target IP: 3.3.3.3  
target port: 80

req 1

source IP: 3.3.3.3  
source port: 80  
target IP: 4.2.1.5  
target port: 267

resp 1



3.3.3.3



10.0.0.2

Clients  
in a local network (LAN)

### Network address translation (NAT)

Private network side		Public network side	
Source IP	Source Port	Target IP	NAT Port
10.0.0.1	837	3.3.3.3	267
10.0.0.2	932	3.3.3.3	413

Remote  
server

### NAT router

1. changes packet's source address from 10.0.0.2:932 to 4.2.1.5:413
2. updates table



837

10.0.0.1



10.0.0.2

932

source IP: 10.0.0.2  
source port: 932  
target IP: 3.3.3.3  
target port: 80

req 2

resp 2

source IP: 3.3.3.3  
source port: 80  
target IP: 10.0.0.2  
target port: 932

Default gateway  
IP address:  
10.0.0.9

Public  
IP address:  
4.2.1.5

Gateway-router  
with  
NAT


source IP: 4.2.1.5  
source port: 413  
target IP: 3.3.3.3  
target port: 80

req 2

resp 2

source IP: 3.3.3.3  
source port: 80  
target IP: 4.2.1.5  
target port: 413

80



3.3.3.3

# NAT: network address translation

- all devices in local network have 32-bit addresses in a “private” IP address space (10.0.0.0/8, 172.16.0.0/12, 192.168.0.0/16 prefixes) that can only be used in local network
- advantages:
  - just **one** IP address needed from provider ISP for *all* devices
  - can change addresses of host in local network without notifying outside world
  - can change ISP without changing addresses of devices in local network
  - security: devices inside local net not directly addressable, visible by outside world

# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- data plane
- control plane

## 4.2 What's inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 The Internet Protocol (IP)

- datagram format
- addressing
- NAT: network address translation
- IPv6

## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action

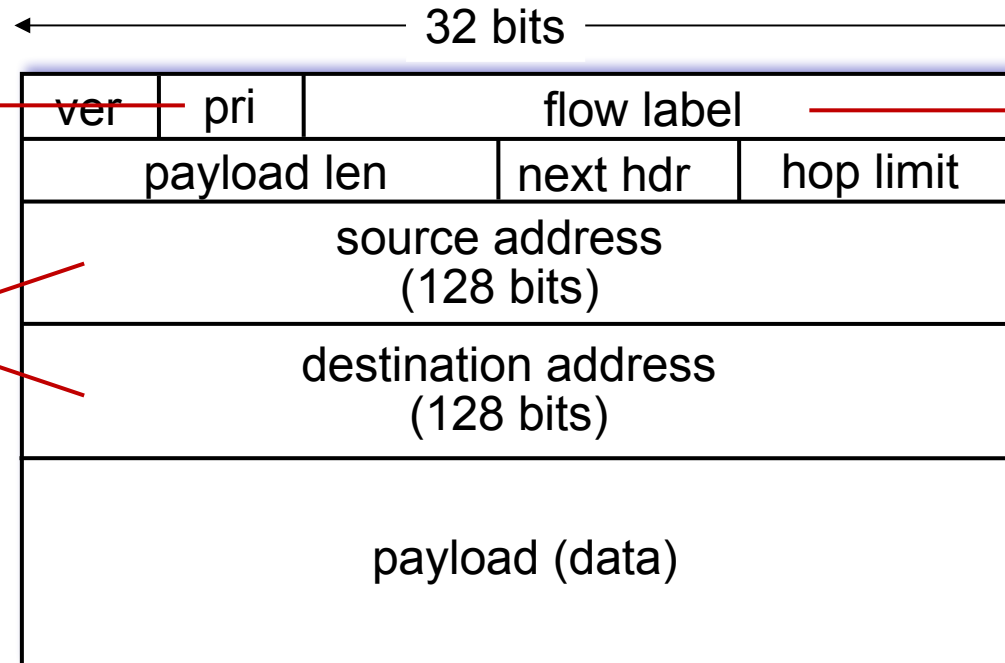
# IPv6: motivation

- **initial motivation:** 32-bit IPv4 address space would be completely allocated
- additional motivation:
  - speed processing/forwarding: 40-byte fixed length header
  - enable different network-layer treatment of “flows”

# IPv6 datagram format

**priority:** identify priority among datagrams in flow

**128-bit** IPv6 addresses




**flow label:** identify datagrams in same "flow." (concept of "flow" not well defined).

What's missing (compared with IPv4):

- no checksum (to speed processing at routers)
- no fragmentation/reassembly
- no options (available as upper-layer, next-header protocol at router)



# Tunneling




# IPv6: adoption

- Google<sup>1</sup>: ~ 40% of clients access services via IPv6 (2023)
- NIST: 1/3 of all US government domains are IPv6 capable

## IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.

Native: 43.36% 6to4/Teredo: 0.00% Total IPv6: 43.36% | Feb 25, 2024



<https://www.google.com/intl/en/ipv6/statistics.html>

# Network layer: “data plane” roadmap

## 4.1 Network layer: overview

- data plane
- control plane

## 4.2 What's inside a router

- input ports, forwarding,
- switching, output ports, scheduling

## 4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6




## 4.4 Generalized Forwarding, SDN

- match + action
- OpenFlow: match + action in action

# Generalized forwarding: match plus action


*Review:* each router contains a **forwarding table** (aka: **flow table**)

- “**match plus action**” abstraction: match bits in arriving packet, take action
  - *destination-based forwarding*: forward based on dest. IP address
  - *generalized forwarding*:
    - many header fields can determine action
    - many action possible:  
drop/copy/modify/log packet




# Flow table abstraction

- **flow**: defined by header field values (in link-, network-, transport-layer fields)
- **generalized forwarding**: simple packet-handling rules
  - **match**: pattern values in packet header fields
  - **actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **priority**: disambiguate overlapping patterns
  - **counters**: #bytes and #packets



# OpenFlow: flow table entries



# OpenFlow: examples

## Destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	51.6.0.8	*	*	*	*	port6

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

## Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	*	*	*	*	22	drop

Block (do not forward) all datagrams destined to TCP port 22 (ssh port #)

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	128.119.1.1	*	*	*	*	*	drop

Block (do not forward) all datagrams sent by host 128.119.1.1

# OpenFlow: examples

Layer 2 destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	22:A7:23: 11:E1:02	*	*	*	*	*	*	*	*	*	port3

layer 2 frames with destination MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3



# OpenFlow abstraction

- **match+action:** abstraction unifies different kinds of devices

## Router

- *match:* longest destination IP prefix
- *action:* forward out a link

## Switch

- *match:* destination MAC address
- *action:* forward or flood

## Firewall

- *match:* IP addresses and TCP/UDP port numbers
- *action:* permit or deny

## NAT

- *match:* IP address and port
- *action:* rewrite address and port

# Chapter 4: done!

- Network layer: overview
- What's inside a router
- IP: the Internet Protocol
- Generalized Forwarding, SDN



*Question:* how are forwarding tables (destination-based forwarding) or flow tables (generalized forwarding) computed?

*Answer:* by the control plane (next chapter)